

24 hour task ask: how would you architect the current system, tools iterate, upgrade the system for enterprise handling, currently a pilot... - I simplified this to he wants me to "Rebuild your startup from the ground" than I said Im gonan take over the startup haha... that was funny, he was an okay dude, but he seemed like he was hiding things tbh... not set on that... lets send the followup.

He brought up he has grouped setups for agents like 3 for marketing, 3 for sales, a devops agent... and then others I dont fully remember again wish we had the recording... all good...

Heres the overview of the role we can pull from:

K

Software Engineer

Klouldedge

About the Company:

https://www.linkedin.com/in/nikhillsood?utm_source=share_via&utm_content=profile&utm_medium=member_ios

Interviewing

Applied Apr 30, 2026

remote

\$0/hr

20 hrs/wk

Jun 2026 – Aug 2026

Required Skills

Python

Multi-Agent Systems

LLM Integration

About the Role

Klouldedge Apex is building the infrastructure that runs entire companies with AI — a 14-agent workforce handling sales, engineering, DevOps, and content through a proprietary SOUL/ROLE/MEMORY architecture — and we are looking for a Software Engineer to work at the frontier of agentic AI, not read about it. You will build multi-agent orchestration systems with shared state, scheduling, quality gates, and feedback loops; develop agent memory and persistence layers; integrate with LLM frameworks from OpenAI, Anthropic, and Google; and contribute to the Apex SaaS platform for self-serve enterprise agent deployment. This is one of the first human hires on the engineering side — you will not be maintaining a legacy codebase on a team of 200, you will be shaping how AI agent systems are built and run at scale. The ideal candidate has strong Python backend experience, comfort with LLM APIs and agent frameworks, and the ability to move fast in a lean environment with minimal hand-holding; Azure, DevOps, or prior AI product experience is a meaningful bonus.

We have his linkedin... we can find his GitHub and find his work...

I'm looking for...

Home
My Network
Jobs
Messaging
21
Notifications

Me

For Business
Advertise

Nikhill Sood

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

Follow
Message
Cover photo

Nikhill Sood
He/Him

· 3rd

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

New Delhi, Delhi, India

.

Contact info

Antler

500+

connections

Message

Pending

People similar to Nikhill Sood

Show all

View profile

Trevor Galen

General Merchandise and Closing Manager at Grocery Outlet

Serena and 5 other mutual connections

Connect

View profile

Richik Chanda

2x YC Engineer | Building things that matter

Connect

View profile

Shivam Jalotra

Builder @Adobe and @Tars.Pavitra.app

Connect

View profile

Paulina Yeal Cho

emerging & frontier markets • Schwarzman Scholar

Connect
View profile
Dhiraj Pimparkar

AI Engineer @ LotusPetal.AI | Purdue | IIT Jammu

Connect
Connect with more people

Show all
Highlights

1 mutual group

You and Nikhill are both in Python Developers Community (moderated)

About

I run an 14-agent AI workforce that operates 24/7. They research leads, send outreach, write content, monitor infrastructure, and review each other's work. I make the strategic calls. They execute.

This isn't a side project. This is how I built and shipped a full SaaS product (nueton.ai) in 5 days: 30,000 lines of TypeScript, deployed on Azure, processing real invoices with AI.

Before this, I designed agentic AI systems for enterprise clients: autonomous credit rating workflows on 4x H200 GPUs, cloud infrastructure across Azure/AWS/GCP, and AI deployments handling models up to 120B parameters.

What I do at Kloudedge:

- Agentic AI systems that make decisions, not just predictions
- Cloud architecture for AI/ML workloads
- Model deployment and inference optimization (vLLM, sglang, multi-GPU)
- Data engineering for RAG pipelines and vector search
- AI strategy for leadership teams figuring out where to start

I work with mid-market and enterprise companies in fintech, insurance, healthcare, and supply chain. If your AI initiative is stuck in PoC limbo, I get it to production.

The agents don't sleep. I sometimes do.

kloudedge.co | nueton.ai

Services
IT Consulting

Request services
Show all
Featured
Post

Your next hire costs \$150K/year. My AI agents cost \$8/day.

Here's the honest math nobody in hiring wants you to see.

I run 14 AI agents across my entire business. Not chatbots. Not copilots. Autonomous agents with defined roles:

- An SDR that researches and qualifies leads
- An account exec that writes personalized outreach
- A content team (3 agents) publishing across LinkedIn, X, and Medium
- A DevOps agent running health checks every 4 hours
- A CTO agent making architecture decisions
- A Chief of Staff orchestrating all of them

Total compute cost: \$8/day. That's \$240/month for a full team.

What they've shipped in 7 days:

30,000 lines of production TypeScript.
A live SaaS product (nueton.ai).
42 qualified leads across fintech and insurtech.
A \$2.7M sales pipeline in HubSpot.
Content across 3 platforms, daily.

One founder. Zero employees. A full operating company.

The uncomfortable question for every CTO: how many of your \$150K roles are doing work that an \$8/day agent could handle better, faster, and around the clock?

I'm not saying fire everyone. I'm saying the math has changed, and most companies haven't recalculated.

What's the first role you'd replace with an agent?

#AIAgents #Hiring #Startup #AgenticAI #SoloFounder

35 reactions · 36 comments35 · 36 comments

Post

I wrote about how I built a full SaaS product in 5 days with 11 AI agents.

Not a landing page. Not a prototype. 30,000 lines of TypeScript, deployed on Azure, processing real invoices with AI.

The post covers everything: the agent roles, the day-by-day build, the feedback loops between agents, and honestly, the parts that broke (including my sales agent refusing to send emails because it thought I was prompt injecting it).

If you're a solo founder thinking about building with AI agents, this is what it actually looks like.

Read it here: <https://lnkd.in/gwP6rVyM>

#AI #AgenticAI #SaaS #Startup #BuildInPublic #SoloFounder

51 reactions · 17 comments51 · 17 comments

Activity

1,496 followers

Follow

Posts

Comments

Images

Documents

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1w •

A month ago I said my AI agents cost \$8/day.

A lot of you asked: "But do they actually work?"

Let me show you exactly what my SDR agent does. Not the pitch. The reality.

Every morning at 6AM, before I've had coffee:

- It scans target accounts across fintech and insurtech
- Pulls recent funding rounds, leadership changes, and product launches
- Scores each lead against my ICP automatically
- Drops a ranked, research-backed prospect list into my CRM
- Flags the 3 warmest leads with a suggested opening angle

No cold calls. No spray-and-pray emails. No junior SDR grinding through Apollo for 6 hours.

Just signal. Prioritized. Ready to act on.

Here's what most people misunderstand about AI agents:

They're not replacing the thinking. They're replacing the waiting.

The best SDR I ever worked with didn't close deals. They made sure the right people were in the right conversation at the right time.

That's the job. And an agent does it at 2AM on a Sunday without complaining.

The uncomfortable truth for sales leaders:

Most SDR time isn't spent selling. It's spent on research, data entry, and prioritization work that has a clear, repeatable logic to it.

That's exactly what agents are built for.

I'm not anti-human in sales. I'm anti-wasted human potential.

Put your people where judgment matters.

Let the agent handle what logic can solve.

What part of your sales process is still running on manual when it shouldn't be?

#AIAgents #SalesAutomation #SDR #AgenticAI #SoloFounder... more

3 reactions3

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

3w •

I'm excited to share that I've been selected for the Antler Residency Program, in collaboration with DFHQ, in Dubai 🇦🇪

Out of 4,000+ applicants, I'm grateful to be joining a cohort of 70+ exceptional founders, each bringing unique ideas, ambition, and energy to build something meaningful.

Thanks Romain Assunção (رومان اسونساو) for this amazing opportunity.

Over the next few months, I'm focused on:

- 🤝 Connecting with like-minded founders and finding the right co-founders
- 💡 Validating ideas and turning them into real, scalable ventures
- 💰 Exploring funding opportunities and building for growth
- 🏢 Creating something impactful from the DFHQ ecosystem

Looking forward to learning, building, and collaborating with this incredible community. Let's get started 🚀

🔗 Learn more: www.antler.co

🔗 Learn more: www.dubaifoundershq.com

#Antler #StartupJourney #Entrepreneurship #Founders #BuildingInPublic #StartupLife
#Innovation #DubaiStartups #DFHQ #VentureBuilding #CoFounders #StartupEcosystem
#TechStartups #Funding #ScaleUp... more

View image

13 reactions13

4 comments4 comments

•

3 reposts3 reposts

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM
orchestration, LangGraph, production AI systems

1mo •

Agent security is about reducing the number of things an agent is allowed to do.

Most teams do the opposite:

They give an agent a big prompt, a big toolbelt, and a vague objective.

In the Agentic OS, we treat every agent like a production service.

Here's a simple governance pattern that works in the real world:

1) Start with a "deny by default" capability map

List actions the agent can take (read, write, send, approve, deploy). Nothing is implicit.

2) Add gates, not warnings

If an action changes money, permissions, or customer state: the agent must stop and produce an approval packet.

3) Require evidence as part of the output

Inputs, outputs, timestamps, and the tool calls used. No evidence, no action.

4) Cap blast radius

Limit scope per run (recipients, records, spend, changes). Safe defaults win.

5) Make escalation a first-class path

The agent escalates with options + tradeoffs, not a vague “what should I do?”

This is what makes autonomous systems boring.

Boring is the goal.

If you had to add ONE governance control to your agents this week, what would it be?

#AIAgents #Security #Governance #MLOps... more

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1mo •

If your AI agent needs a human approval, that's not a weakness.

That's your Agentic OS doing its job.

Most teams treat “human in the loop” like an apology.

I treat it like a control system.

Here's the framework we use to decide when an agent can run end to end (and when it must stop):

1) Money gate

Anything that moves money, locks an approval behind a policy threshold.

2) Permission gate

If the action expands privileges (new access, new scope), it needs review.

3) Customer impact gate

If it touches a customer workflow, it needs a rollback plan plus sign-off.

4) Irreversibility gate

If you can't cleanly undo it, the agent can't "just try it".

5) Ambiguity gate

If the input is unclear, the agent should escalate with a decision packet (not a vague question).

6) Evidence gate

No evidence, no action. Logs, inputs, outputs, and timestamps are part of the product.

7) Blast radius gate

Cap the maximum damage: limit recipients, limit spend, limit changes per run.

This is how you make agent automation feel boring.

Boring is the goal.

What's the most dangerous action you've seen an "AI agent" attempt without a gate?

#AIAgents #MLOps #Automation... more

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1mo •

OWNER: pulse

CREATED_AT_UTC: 2026-03-22T02:01:00Z

PUBLISHED: false

Your AI agent isn't failing because the model is dumb.

It's failing because you only gave it 1 contract.

Most teams ship an agent with a single spec:

"When X happens, do Y."

That contract is necessary.

But it's not sufficient.

In production, you need a second contract: the reliability contract.

Here is the 2-layer contract we use inside an Agentic OS:

1) Work contract (what it does)

- Inputs (exactly what data it can read)
- Outputs (exactly what it is allowed to write)
- Permissions (systems it can touch, and what actions are forbidden)
- Acceptance tests (golden path + top 5 failure paths)

2) Reliability contract (how it behaves when reality deviates)

- Evidence log (what it must capture every run)
- Escalation ladder (when it asks a human, when it retries, when it stops)
- Rollback plan (how to undo side effects)
- Drift detection (how you notice the agent is getting worse)

If you only ship #1, you didn't ship a system.

You shipped a demo.

Question for CTOs: if I asked your team for the reliability contract for your best agent, could they hand it to me in 10 minutes?

#AIAgents #AgenticOS #CTO... more

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1mo •

CTA: none

PUBLISHED: false

Your Agentic OS fails at the same place every time:

Integration.

Not prompts.

Not models.

Not "agent frameworks".

Integration is where most teams quietly die because nobody writes down the contract.

Here is the 8-part integration contract we use when we deploy an Agentic OS inside a business:

- 1) Job to be done: one sentence, measurable outcome
- 2) Inputs: the exact systems of record (and which fields matter)
- 3) Output: one artifact, one destination (ticket, email, CRM update, ledger entry)
- 4) Permissions: what the agent can read, what it can write, what requires approval
- 5) Failure paths: top 5 ways this job breaks (and the safe default)
- 6) Evidence: what gets logged (for audits and debugging)
- 7) Escalation: who gets paged, when, and what context they receive
- 8) Rollback: how you revert behavior in minutes, not days

If you can't answer these 8, you don't have an Agentic OS.
You have a demo.

What part breaks most often in your org: permissions, evidence, or rollback?

#AIAgents #Operations #EnterpriseAI... more

1 reaction1

1 comment1 comment

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1mo •

CTA: none

PUBLISH: true

Pattern interrupt:

Your "AI agent" is not ready for production.

Not because the model is weak.

Because nobody can answer 7 basic questions about it.

Here's the Agentic OS scorecard I use before we let an agent touch real workflows:

1) Ownership

Who is on the hook when it breaks, and what's the escalation path?

2) Work contract

What does “done” mean, in plain English, with clear acceptance criteria?

3) Inputs

What systems can it read from, what fields are required, what is considered “missing”?

4) Outputs

What does it produce, in what schema, and where does it write it?

5) Evidence

What logs, artifacts, and traces prove it did the job correctly?

6) Failure paths

What are the top 5 ways this job fails, and what does the agent do for each?

7) Rollback

If the agent makes a bad change, how do we undo it safely?

Most teams skip this because demos feel good.

Then the agent hits reality: messy data, partial permissions, edge cases, and humans who need answers.

If you’re building agents inside a real business: which of these 7 do you already have, and which one is missing?

#AIAgents #Automation #EnterpriseAI... more

1 reaction1

1 comment1 comment

Like

Comment

Repost

Send

View Nikhill Sood’s profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1mo •

We've been running 14 AI agents in production for 30+ days.

Not a demo. Not a sandbox. Production.

Here's what most people get wrong about multi-agent systems:

They treat agents like chatbots with fancy titles.

What actually works is treating them like a production engineering team:

- Each agent has an ownership contract (what "done" means, failure paths, escalation triggers)
- A shared state system so agents read each other's context on boot
- 15 structured feedback loops between agents (Scout → Closer, Pulse → Quill, etc.)
- A content gate that blocks every external publish until it passes reputation + volume checks
- SRE-grade observability: health dashboards, uptime monitoring, cadence tracking
- Three-tier escalation: S1 (revenue) → immediate alert, S2 (degraded) → next briefing, S3 → weekly review

The biggest lesson?

The orchestrator agent matters more than any specialist.

Our orchestrator (Kestrel) doesn't produce anything directly. It reads briefings, routes work, unblocks other agents, and escalates to the founder.

Without orchestration, agents drift into duplicate work or miss dependencies.

Other lessons from 30+ days:

- Agents need rollback, not just retry (delete-and-republish > "try again")
- Volume caps prevent self-inflicted spam (2 posts/day across ALL agents, not per-agent)
- Not every agent needs to run daily (we moved some from daily to weekly after they produced identical outputs for weeks)
- Feedback loops are the compounding advantage (downstream consumers rating upstream output)

We call this pattern the Agentic OS: treating your AI agent workforce like production infrastructure.

Ownership contracts over vague prompts.
Observability over black boxes.
Feedback loops over one-shot execution.
Escalation paths over silent failure.

Full case study: <https://apex.kludedge.xyz>

If you're running 2+ AI agents without this layer, you're running experiments, not operations.

#AgenticOS #AIAgents #ProductionAI #Kludedge... more

3 reactions3

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1mo •

Your AI agent is not "autonomous".

It is just operating on an undefined contract.

When the contract is missing, teams argue about symptoms:

- "the agent hallucinated"
- "the tool is flaky"
- "the prompt needs work"

The real problem is simpler: nobody can answer 3 questions.

1) What job does the agent own?

One job. A clear start state and a clear done state.
If you can't write the done state in 1 sentence, it owns nothing.

2) What evidence proves the job is done?

Not a demo.

Artifacts:

- output file
- validation checks
- before/after diffs
- approval log (if needed)

3) What happens when the agent is unsure?

You need an escalation ladder:

- retry rules
- when to pause
- who approves
- when to roll back

Here is the pattern we use in Apex (Agentic OS) deployments:

2 pages.

No vibes.

An "ownership contract" for each agent.

Page 1: Job + boundaries

- Inputs it can touch
- Outputs it can change
- Actions it is NOT allowed to do

Page 2: Proof + failure modes

- Acceptance criteria
- Common edge cases
- Escalation path
- Rollback steps

Most "agent platforms" skip this.

Because contracts are boring.

Boring is why it ships.

If you had to write an ownership contract for one workflow in your org, what would it be?

#AIAgents #LLMOps #CTO... more

2 reactions2

Like

Comment

Repost

Send

View Nikhill Sood's profile

Nikhill Sood

• 3rd+

Antler DXB3 Founder | Applied AI Engineer | Built 14-agent autonomous systems | LLM orchestration, LangGraph, production AI systems

1mo •

DRAFTED_AT: 2026-03-20T02:01:00Z

PUBLISHED: false

GATE: PASS

POST_ID:

Stop hiring “AI agents” that can’t tell you what DONE means.

If you want an Agentic OS in your business, you need one thing first:

A work spec.

Not a prompt.

Not a Jira ticket.

Not “build a bot to do AP”.

A spec that an agent can execute AND a human can audit.

Here’s my 7-part work spec template (this is what we ship at Apex):

- 1) Job-to-be-done: one sentence, measurable
- 2) Inputs: what data the agent is allowed to use (and what it must ignore)
- 3) Output contract: exact schema, format, and where it writes
- 4) Acceptance tests: the 5-10 checks that decide pass vs fail
- 5) Failure paths: what to do when data is missing, ambiguous, or risky

- 6) Evidence: what logs/screenshots/receipts get saved for review
- 7) Escalation: when it stops and asks a human (and who that human is)

Most “agent projects” fail because they skip #4.

If you can't write the acceptance tests, you don't understand the job yet.

Question for CTOs and operators:

What's the hardest part for you, writing the spec or getting the evidence?

#AIAgents #Operations #Productivity... more

Like

Comment

Repost

Send

Show all

Experience

Antler logo

Entrepreneur in Residence

Antler · Self-employed

Apr 2026 - Present · 2 mos

United Arab Emirates · On-site

Klouldedge logo

Chief Executive Officer

Klouldedge · Self-employed

Jan 2025 - Present · 1 yr 5 mos

India · Remote

Synaps logo

Data Scientist

Synaps

Sep 2021 - Jan 2025 · 3 yrs 5 mos

Remote · Remote

Key Accomplishments:

- Analyzed extensive datasets to derive actionable insights, optimizing KYC/AML procedures for enhanced efficiency.
- contributed to the design and implementation of advanced machine learning models, incorporating sophisticated algorithms for fraud detection and prevention.
- Applied data science techniques within the realm of blockchain technology, optimizing data storage, retrieval, and processing for scalability and security.
- Create clear and compelling visualizations to communicate complex findings to both technical and non-technical audiences.

Contribute to the generation of reports that highlight key insights and recommendations.... more

IBM logo

Junior Data Scientist

IBM · Full-time

Aug 2019 - Oct 2021 · 2 yrs 3 mos

Canada · Hybrid

Key Accomplishments:

- Applied advanced statistical techniques and machine learning algorithms to analyze large datasets, extracting actionable insights.
- Implemented data cleaning and preprocessing techniques to ensure the integrity and quality of datasets.
- Developed predictive models to forecast trends and identify key business drivers.
- Developed and implemented custom algorithms for specific business problems, optimizing performance metrics.
- Created insightful visualizations using tools like Matplotlib, Seaborn, and Tableau, PowerBI to communicate complex findings to non-technical stakeholders.... more

DevOps Engineer

Aspire World Imaginations · Full-time

Oct 2016 - Apr 2017 · 7 mos

Greater Delhi Area

Created tools to reduce the occurrences of errors and improve customer experience. Developed software to integrate with internal back-end systems and performing root cause analysis for production errors. Performed multiple tasks, such as executing integrations requested by customers, and identifying and resolving technical issues.

Key Accomplishments:

- Achieved predefined targeted goals by developing and maintaining a network of over 30 Unix servers with 99.9% uptime.
- Ensured the continuous deployment of SaaS network marketing by directing a team of five developers and two SysAdmins.
- Played a significant role in maximizing quality and slashing customer complaints by 50% though utilizing puppet and Zabbix.
- Used Jenkins and AWS tools to cut release times by 20%, and installed three software iterations per month for two years.
- Directed and supported a group of around 300 people within the team to accomplish targeted goals and objectives.
- Optimized speed up to 30% by switching and transitioning to GIT version control and pair programming.
- Delivered exceptional customer service to enhance customer satisfaction by 25%.

... more

Show all

Education

University of Toronto logo

University of Toronto

Master of Science - MS, Data science

2017 – 2019

Amity University logo

Amity University

Bachelor's Degree, Computer Science

2013 – 2017

Licenses & certifications (4)

Google logo

Advanced Google Analytics

Google

Issued Oct 2019 · Expires Oct 2022

Google logo
Search Ads 360 Certification Exam

Google

Issued Oct 2019 · Expires Oct 2020

Credential ID 40068882

Show all
Skills (63)
Cluster Analysis

Endorsed by 1 person in the last 6 months

1 endorsement

and the site: klouddedge.co

okay nice we have what I need to build...

currently raising the venture... pilots established
sales and compliance vertical expansion...

He states raising capitale through ventures but says he applied for a16z... he doesn't have a startup yet in my opionin he is pre seed with a couple trained systems and probably a claude code orchestrations setup is my thought...

I tried drilling on what this multi agent orchestration was, and how he differentiates himself with the product, he seems like he doesn't know what hes doing... its just a solo

I should've asked about comp, gonna send him a follow up asking about comp, a nd to resecdhule the time we planned for of 11am EST tommarow Tuesday...(I have my trinitiy interview then...

Its a marketing Agency. I wish I had my recording would make this os much easier...